

$$+ \text{Code} \big) - \big(+ \text{Text}$$

```
load_iris()
```

```
datasets=pd.DataFrame(df.data)
print(datasets)
```

	0	1	2	3
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

```

145  6.7  3.0  5.2  2.3
146  6.3  2.5  5.0  1.9
147  6.5  3.0  5.2  2.0
148  6.2  3.4  5.4  2.3
149  5.9  3.0  5.1  1.8

```

[150 rows x 4 columns]

```
datasets.columns=df.feature_names
```

```
datasets.head()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

Next steps: [Generate code with datasets](#) [View recommended plots](#) [New interactive sheet](#)

```

## here i will keep my independent features X and dependent features as y
X=datasets
y=df.target

```

y

```

array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2])

```

```
## train test split we do here
```

```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=42)

```

X_train

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	
81	5.5	2.4	3.7	1.0	
133	6.3	2.8	5.1	1.5	
137	6.4	3.1	5.5	1.8	
75	6.6	3.0	4.4	1.4	
109	7.2	3.6	6.1	2.5	
...	
71	6.1	2.8	4.0	1.3	
106	4.9	2.5	4.5	1.7	
14	5.8	4.0	1.2	0.2	
92	5.8	2.6	4.0	1.2	
102	7.1	3.0	5.9	2.1	

105 rows x 4 columns

Next steps: [Generate code with X_train](#) [View recommended plots](#) [New interactive sheet](#)

```

# standardizing the dataset by this output will be correct
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()

```

```
## scaler.inverse_transform(X_train) we can get back from standardizing
```

[-0.29283662,	-0.74242333,	0.19235097,	0.08245999],
[-0.89573553,	0.93659559,	-1.38370397,	-1.40568508],
[-0.17225683,	-0.02284379,	0.19235097,	-0.05282593],
[2.23933883,	1.89603497,	1.65166111,	1.30003323],
[-1.49863445,	0.4568759 ,	-1.44207638,	-1.40568508],
[0.43064208,	-0.26270364,	0.25072338,	0.08245999],
[-0.17225683,	-1.22214302,	0.65933022,	1.0294614],
[-0.4134164 ,	2.85547435,	-1.44207638,	-1.40568508],
[0.18948252,	-0.02284379,	0.54258541,	0.75888956],
[-0.05167705,	-0.74242333,	0.71770262,	0.89417548],
[0.18948252,	-1.94172256,	0.07560616,	-0.32339776],
[-0.53399618,	-0.02284379,	0.36746819,	0.35303182],
[0.43064208,	0.93659559,	0.89281984,	1.43531914],
[-0.4134164 ,	-1.70186271,	0.07560616,	0.08245999],
[-0.53399618,	2.13589481,	-1.26695916,	-1.13511325],
[-1.01631531,	-1.70186271,	-0.33300068,	-0.32339776],
[0.67180165,	-0.74242333,	0.83444743,	0.89417548],
[-1.01631531,	0.69673574,	-1.44207638,	-1.40568508],
[-1.01631531,	0.4568759 ,	-1.55882119,	-1.40568508],
[-0.4134164 ,	-1.46200287,	-0.04113865,	-0.18811184],
[1.033541 ,	-0.02284379,	0.65933022,	0.62360365],
[-1.1368951 ,	0.21701605,	-1.38370397,	-1.40568508],
[-0.05167705,	-0.50256349,	0.71770262,	1.57060506],
[-1.01631531,	0.93659559,	-1.38370397,	-1.40568508],
[-1.01631531,	1.17645543,	-1.32533157,	-0.86454142],
[0.06890273,	0.4568759 ,	0.54258541,	0.75888956],
[-0.89573553,	-1.22214302,	-0.50811789,	-0.18811184],
[1.27470056,	0.4568759 ,	1.06793706,	1.43531914],
[0.18948252,	-0.74242333,	0.71770262,	0.48831773],
[0.3100623 ,	-0.98228318,	1.00956465,	0.2177459],
[2.23933883,	-0.02284379,	1.30142668,	1.43531914],
[-0.4134164 ,	-1.22214302,	0.07560616,	0.08245999],
[-1.73979401,	-0.26270364,	-1.44207638,	-1.40568508],
[-1.8603738 ,	-0.02284379,	-1.6171936 ,	-1.540971],
[0.18948252,	-1.94172256,	0.65933022,	0.35303182],
[1.63643991,	0.4568759 ,	1.24305427,	0.75888956],
[-1.49863445,	0.21701605,	-1.38370397,	-1.40568508],
[-0.89573553,	1.17645543,	-1.44207638,	-1.27039917],
[-1.73979401,	-0.02284379,	-1.50044878,	-1.40568508],
[0.55122187,	-1.22214302,	0.60095781,	0.35303182],
[0.55122187,	0.93659559,	1.00956465,	1.57060506],
[-1.49863445,	0.93659559,	-1.44207638,	-1.27039917],
[1.15412078,	-0.02284379,	0.95119225,	1.16474731],
[0.55122187,	0.69673574,	1.24305427,	1.70589097],
[-1.37805466,	0.4568759 ,	-1.50044878,	-1.40568508],
[0.3100623 ,	-0.26270364,	0.484213 ,	0.2177459],
[0.79238143,	-0.50256349,	0.4258406 ,	0.35303182],
[0.43064208,	-0.50256349,	0.54258541,	0.75888956],
[1.39528035,	0.4568759 ,	0.484213 ,	0.2177459],
[0.67180165,	0.4568759 ,	0.83444743,	1.43531914],
[-0.89573553,	1.89603497,	-1.32533157,	-1.40568508],
[1.27470056,	0.21701605,	0.89281984,	1.16474731],
[0.06890273,	-0.02284379,	0.19235097,	0.35303182],
[0.79238143,	-0.02284379,	0.77607503,	1.0294614],
[-0.17225683,	-0.98228318,	-0.21625586,	-0.32339776],
[-0.77515575,	0.74242333,	-0.01723376,	0.2177459],
[0.3100623 ,	-0.02284379,	0.4258406 ,	0.2177459],
[-1.61921423,	-1.70186271,	-1.50044878,	-1.27039917],

```
## we implement here the linear regression
## we import the linear regression from sklearn
from sklearn.linear_model import LinearRegression
# cross validation
from sklearn.model_selection import cross_val_score
```

```
regression=LinearRegression()  
regression.fit(X_train,y_train)
```

LinearRegression()

```
mse=cross_val_score(regression,X_train,y_train,scoring='neg mean squared error',cv=5)
```

```
np.mean(mse)
```

```
np.float64(-0.05617171940989511)
```

```
#prediction
```

```
reg_predict=regression.predict(X_test)
```

```
reg_predict
```

```
array([[ 1.24069097, -0.04537609,  2.24501083,  1.35143666,  1.29775083,
         0.01024241,  1.05031173,  1.82525399,  1.37084413,  1.06699186,
        1.70363485, -0.08712067, -0.165166   , -0.07724353, -0.03380619,
        1.40167699,  2.00651252,  1.04725931,  1.28368327,  1.97600474,
        0.01782354,  1.59952875,  0.079732   ,  1.92307532,  1.8621986   ,
        1.8790815   ,  1.80251247,  2.04196713,  0.01873817,  0.01291496,
       -0.15365607, -0.08046738,  1.18506728, -0.00461982, -0.02934265,
        1.68665136,  1.29088786, -0.07995434, -0.09076782, -0.16795331,
        1.75520461,  1.37514144,  1.3174234   , -0.07193336, -0.1131512  ]])
```

```
import seaborn as sns
sns.distplot(reg_predict-y_test)
plt.show()
```


```
<ipython-input-22-0c78bb25d2d9>:2: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(reg_predict-y_test)
```



```
from sklearn.metrics import r2_score
score=r2_score(reg_predict,y_test)
```

```
score
```

```
0.9453000057840795
```

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
```

```
def gradient_descent_visualization(x, y, learning_rate=0.01, iterations=1000):
    m_curr = 0
    b_curr = 0
    n = len(x)
    cost_history = []
```

```
plt.figure(figsize=(10, 5))
```

```
for i in range(iterations):
    y_predicted = m_curr * x + b_curr
    cost = (1/n) * sum((y - y_predicted) ** 2)
    cost_history.append(cost)
```

```

md = (-2/n) * sum(x * (y - y_predicted))
bd = (-2/n) * sum(y - y_predicted)

m_curr = m_curr - learning_rate * md
b_curr = b_curr - learning_rate * bd

# Plot every 100 iterations
if i % 100 == 0:
    print(f"Iteration {i} | Cost: {cost}")
    plt.scatter(x, y, color="red")
    plt.plot(x, y_predicted, label=f"Iteration {i}")
    plt.xlabel("Feature (X)")
    plt.ylabel("Target (Y)")
    plt.title("Gradient Descent Line Updates (Iris Feature)")
    plt.legend()
    plt.pause(0.5)
    plt.clf()

# Final best fit line
y_predicted = m_curr * x + b_curr
plt.scatter(x, y, color="red")
plt.plot(x, y_predicted, color="green", label="Best Fit Line")
plt.xlabel("Feature (X)")
plt.ylabel("Target (Y)")
plt.title("Final Best Fit Line on Iris Data")
plt.legend()
plt.show()

# Plot cost function convergence
plt.figure(figsize=(8, 5))
plt.plot(range(iterations), cost_history, color='blue')
plt.xlabel("Iterations")
plt.ylabel("Cost")
plt.title("Cost Function Convergence")
plt.show()


print(f"Final slope (m): {m_curr}, intercept (b): {b_curr}")

# Load Iris dataset
iris = load_iris()
X = iris.data[:, 0] # sepal length (feature)
Y = iris.data[:, 1] # sepal width (target)


gradient_descent_visualization(X, Y)

```


Iteration 0 | Cost: 9.535999999999994



Iteration 100 | Cost: 0.3877364149031411



Iteration 200 | Cost: 0.37296089599581966



Iteration 300 | Cost: 0.35926812147799686

